

How to Build the Wiki: A Step-by-Step Guide

Companion to “The Pattern Everyone’s Quietly Building”

By Mercedes Perez Capilla

This is the practical companion to the main piece. If you want to understand *why* the LLM Wiki pattern matters, read that one first. This one tells you exactly how to set yours up using Claude Code, following Karpathy’s pattern.

You don’t need to be a developer. If you can create a folder and run a command in a terminal, you can do this. Total setup time: about an hour, plus your first ingest.

Before you start

You’ll need:

1. **Claude Code** installed. If you don’t have it yet, follow the install instructions at docs.claude.com.
2. **Obsidian** (free) for the view layer. Download from obsidian.md.
3. **A folder** somewhere on your machine that will become your wiki. I use `~/wiki`. You can call yours anything.
4. **One source to ingest** — an article, a PDF, a set of notes. Anything you want your wiki to know about.

Optional but recommended: - A git repo for the wiki folder, so you get version history for free.

Step 1 — Create the folder structure

Open a terminal and run:

```
mkdir -p ~/wiki/{raw,context,wiki}
cd ~/wiki
touch CLAUDE.md
touch wiki/index.md wiki/log.md
```

Your folder now looks like this:

```
~/wiki/
├── CLAUDE.md          ← empty for now, will become your schema
├── raw/               ← source documents go here
├── context/           ← your Context Vault lives here
└── wiki/              ← Claude's output lives here
    ├── index.md       ← catalogue of every page
    └── log.md          ← append-only record of operations
```

Step 2 — Read Karpathy's gist

Before you write your `CLAUDE.md`, read Karpathy's original gist:

gist.github.com/karpathy/442a6bf555914893e9891c11519de94f

It's about a 10-minute read. Don't skip it. The whole pattern rests on the three operations he describes — ingest, query, lint — and the three-layer architecture (raw, wiki, schema). Understanding these is what lets you adapt the pattern to your own work later.

Step 3 — Write your CLAUDE.md schema

This is the most important file in the system. It tells Claude the rules of *your* wiki — what to do, what not to do, how to name things, what conventions to follow.

The good news: you don't have to write it from scratch. Open Claude Code in your wiki folder and ask Claude to help. Here's the prompt I used (adapt to your domain):

"I'm setting up an LLM Wiki based on Andrej Karpathy's pattern in this folder. My wiki will cover [your domain — e.g. 'my work in financial services AI', or 'my research into language learning', or 'my personal knowledge base across work, family, and projects'].

Read the Karpathy gist at gist.github.com/karpathy/442a6bf555914893e9891c11519de94f and then help me write a `CLAUDE.md` for this wiki. The schema should cover: — the folder structure and what belongs where — page naming conventions and structure — how to handle entities, concepts, and source summaries — what to do when a new source contradicts an existing page (flag, don't overwrite) — what's private vs shareable

Ask me clarifying questions before writing. Keep it under 200 lines — schemas drift if they get too long.”

Claude will ask you questions about your domain, your conventions, your privacy preferences. Answer them honestly — the schema is only as good as the answers you give. The richer your `CLAUDE.md`, the better the wiki will perform.

Step 4 — Build your Context Vault (if you don't have one)

The Context Vault is what makes Claude *know you* at the start of every session. It lives in `~/wiki/context/`.

Follow Allie K Miller's approach at alliekmiller.com/ai-context-vault — she has a set of interview prompts that produce layered context documents covering who you are, what you work on, how you think, and how you communicate.

Once you've built it, add a line to your `CLAUDE.md` telling Claude to read it at the start of every session:

“At the start of every session, read all files in `~/wiki/context/` to ground yourself in who I am and what I'm working on.”

This single line is the one that makes the AI feel like it actually knows you.

Step 5 — Run your first ingest

Pick one source. Something you've actually read. An article, a paper, your favourite blog post — doesn't matter what.

1. Save it to `~/wiki/raw/` as a markdown file. Name it date-prefixed: `2026-05-26-the-article-title.md`. If it's a PDF, just drop the PDF in.
2. Open Claude Code in your wiki folder.

3. Tell Claude:

“Ingest the file in `raw/` called `[filename]` into the wiki. Follow the rules in `CLAUDE.md`. Before you write anything, show me the takeaways and the wiki pages you plan to create or update, and wait for my confirmation.”

Claude will read the source, surface the key takeaways, and propose which pages to create and update. Review the plan. If it looks right, say “go ahead.” If it looks wrong, push back — that’s how the schema gets refined.

After the ingest:

- Check `~/wiki/wiki/` — there should be new pages.
- Check `~/wiki/wiki/index.md` — it should list them with one-line summaries.
- Check `~/wiki/wiki/log.md` — there should be one new line recording the ingest.

If any of these are missing, your schema needs tightening. Ask Claude to update `CLAUDE.md` to enforce them on every ingest.

Step 6 — Open Obsidian on the same folder

In Obsidian: 1. *Open folder as vault* 2. Point it at `~/wiki/` 3. Open graph view (Ctrl/Cmd + G)

You should see your wiki as a graph — pages as nodes, `[[wikilinks]]` as edges. Right now it’s tiny. Ingest five more sources and it starts to look like a knowledge graph. Ingest fifty and it starts to look like a brain.

Step 7 — Repeat

That’s the whole loop:

- Drop a source into `raw/`
- Run an ingest
- Watch the wiki grow
- Periodically run a lint check (ask Claude: *“lint the wiki — find contradictions, orphans, stale claims”*)

After a week, your wiki will be a few dozen pages and you'll start to notice it answering questions in your own voice, citing things you've forgotten you read, and surfacing connections you didn't know existed. That's the compounding effect.

Common questions

How long until it's useful? Usefully personalised after one Context Vault and three or four ingests. Genuinely valuable after fifteen to twenty ingests. Hard to give up after fifty.

What should I ingest? Start with things you've already read and want to remember. Articles you saved months ago. Books you took notes on. Meeting notes from projects that still matter. Don't ingest everything you encounter — be selective. The wiki rewards quality over quantity.

Do I need to back it up? Yes. The wiki is just text files, so `git init` in the wiki folder gives you full version history. Push to a private GitHub repo for off-machine backup.

What about live sources like Gmail or Slack? You can connect MCP connectors to give Claude access to these, but don't auto-ingest everything — the noise will swamp the signal. Instead: weekly, ask Claude to surface the handful of items worth keeping. You confirm, Claude saves the curated set to `raw/inbox/`, then runs the normal ingest on them. The wiki doesn't know or care where a source came from — it just sees markdown files in `raw/`.

What about work data? If you work in a regulated industry, do not put work data in a personal wiki without your organisation's approval. Build a separate wiki for that, on approved infrastructure, with appropriate governance. The pattern works the same; the boundaries don't.

Where to go next

Once the basic loop is running, you'll start to notice operations you repeat. Turn them into slash commands in `.claude/commands/`. Four good ones to start with:

- `/ingest <source>` — automate the workflow above
- `/lint` — periodic health check
- `/save-to-vault` — capture session insights before closing
- `/daily` — morning digest of what changed and what needs attention

I'll cover these in more depth in Part 2 of the main series, along with what other practitioners in the RepresentAI community are building.

Credits

This guide combines two patterns:

- **Andrej Karpathy's LLM Wiki** for the architecture and operations:
gist.github.com/karpathy/442a6bf555914893e9891c11519de94f
- **Allie K Miller's AI Context Vault** for the layered personal context:
alliekmiller.com/ai-context-vault

If you build something interesting on top, share it — that's how the pattern keeps evolving.

Mercedes Perez Capilla is a Product Lead for Transformation and AI Adoption in financial services. She is a Claude Certified Architect candidate and part of the Financial Services group at RepresentAI.